

Octopus: A Secure and Anonymous DHT Lookup

Qiyang Wang

Department of Computer Science
University of Illinois at Urbana-Champaign
IL, U.S.A
qwang26@illinois.edu

Nikita Borisov

Department of Electrical and Computer Engineering
University of Illinois at Urbana-Champaign
IL, U.S.A.
nikita@illinois.edu

Abstract—Distributed Hash Table (DHT) lookup is a core technique in structured peer-to-peer (P2P) networks. Its decentralized nature introduces security and privacy vulnerabilities for applications built on top of them; we thus set out to design a lookup mechanism achieving both security and anonymity, heretofore an open problem. We present the design of Octopus, which uses attacker identification mechanisms to discover and remove malicious nodes, severely limiting an adversary's ability to carry out active attacks, and splits lookup queries over separate anonymous paths and introduces dummy queries to achieve high levels of anonymity. We analyze the security of Octopus by developing an event-based simulator to show that the attacker discovery mechanisms can rapidly identify malicious nodes with low error rate. We calculate the anonymity of Octopus using probabilistic modeling and show that Octopus can achieve near-optimal anonymity. We evaluate Octopus's efficiency on Planetlab and show that Octopus has reasonable lookup latency and low bandwidth overhead.

Keywords—Anonymity, security, DHT, lookup

I. INTRODUCTION

Structured peer-to-peer networks, such as Chord [1] or Kademlia [2], allow the creation of scalable distributed applications that can support millions of users. They have been used to build a number of successful applications, including P2P file sharing (Overnet, Kad, Vuze DHT¹) and content distribution (CoralCDN [3]), and many others have been proposed, such as distributed file systems [4], anonymous communication systems [5]–[9], and online social networks [10], [11]. At the heart of these networks lies a distributed hash table (DHT) lookup mechanism that implements a decentralized key–value store. The DHT allows efficient storage and coordination among very large collections of nodes; however, its decentralized nature creates a number of security and privacy vulnerabilities. Because peers have to rely on other peers to determine the state of the network, malicious nodes could provide misinformation to misdirect an honest user's lookups [12]. Likewise, nodes can profile the lookup activities of other nodes and learn what files or websites they are interested in or who their friends may be. Recent research shows that anonymous communication systems based on anonymity-deficient DHT lookups have severe vulnerabilities to information leak attacks [13], [14].

¹<http://www.vuze.com/>

To address these issues, our goal is to design a lookup mechanism that achieves both security and anonymity (where anonymity means no information is revealed about which nodes are looking up which values), heretofore an open problem. We note that security is a *necessary condition* for anonymity, because without security, malicious nodes can misdirect the lookup towards colluding nodes and learn about the lookup target. On the other hand, security is *not a sufficient condition* for anonymity. Some existing lookup schemes designed to resist active attacks are not suitable to build anonymous DHT systems, e.g., due to heavily relying on redundant transmission that leaks information about the lookup initiator and/or target [15]–[17]. Furthermore, even with a secure lookup that itself does not cause information leak, an inappropriate design of the DHT system can still lead to anonymity vulnerabilities [14].

Our contributions in this work include:

1) We propose a suite of novel security mechanisms—*attacker discovery* for DHT systems. Our mechanisms proactively identify and remove malicious peers. We develop an event-based simulator to show that our identification mechanism is capable of rapidly discovering malicious nodes with low error rate. For a network with 20% malicious nodes, it can correctly identify all attacking nodes within 30 minutes. We compare our scheme with Halo [16], a state-of-the-art secure DHT system, and show that our scheme provides better robustness against active attacks. Furthermore, our design does not rely on redundant transmission, and thus is suitable to construct anonymous DHT systems.

2) With the proposed security mechanisms, we put forward a secure and anonymous DHT lookup *Octopus*. Octopus splits individual queries used in a lookup over multiple anonymous paths, and introduces dummy queries, to make it difficult for an adversary to learn the eventual target of a lookup. We use probabilistic modeling with the help of simulation to calculate the information leak, so that users can know how much anonymity can be actually provided by the system. We show that Octopus provides near-optimal anonymity for both the lookup initiator and target. In a network of 100 000 nodes with 20% malicious nodes, Octopus only leaks 0.57 bit of information about the initiator and 0.82 bit of information about the target; these are 4–6 times

better than what previous work [7], [8] was able to achieve.

3) For performance evaluation, we measure the lookup latency of Octopus on Planetlab with 207 nodes, and compare it with the base-line scheme Chord [1] and Halo [16]. The lookup latency of Octopus is comparable to that of Chord, and even better than that of Halo. While Octopus incurs relatively higher communication overhead than Chord and Halo to provide extra security and/or anonymity guarantees, the bandwidth consumption of Octopus is still manageable at only few Kbps for each node.

The remainder paper is organized as follows. Section II presents the system model. We describe our security and anonymity mechanisms in Sections III and IV, respectively. The efficiency evaluation is provided in Section V and Section VI presents the related work. We conclude in Section VII.

II. SYSTEM MODEL

A. Threat Model

In the same vein as previous work [5]–[9], [17], [18], we do not consider a global adversary that is capable of controlling the whole network and observing all communication traffic. Such a global adversary seems impractical in large-scaled P2P networks. Instead, we assume a partial adversary that controls a fraction f of all nodes in the network (f is typically assumed to be up to 20%). Malicious nodes can behave in an arbitrarily malicious way, such as intercepting, modifying, or dropping any messages going through them, or injecting fake messages to any other nodes. We also assume that malicious nodes can log any messages they have seen and have access to a high-speed communication channel to share any information with very low transmission delay.

Also similar to previous work, we do not attempt to solve the problem of the Sybil attack [19] in this work. Defending from a Sybil attack is an interesting research area that has drawn a lot of attentions; a number of effective solutions have been proposed, such as [20]–[22] and may be possible to integrate into the Octopus design as an extension.

B. Design Goals

The major goal for security is to avoid lookups being biased by malicious nodes. In other words, given a lookup target, it should not be possible to misdirect the lookup path or bias the final lookup result.

Pfitzmann and Hansen defined several relevant anonymity properties for message-based communication, such as sender and receiver anonymity [23]. We consider equivalent properties in the context of DHT lookups.

- *Initiator anonymity*: given a lookup target, it should not be possible to determine its initiator.
- *Target anonymity*: given a lookup initiator, it should not be possible to determine its target.

- *Query unlinkability*: given several queries with known targets, it should not be possible to find out if they came from the same initiator.

III. SECURITY MECHANISMS OF OCTOPUS

A. Problem Description

In a DHT system, like Chord [1], each node is assigned a unique ID associated with its IP address, and owns the IDs from itself to its direct predecessor on the ring. Each node X maintains a list of $\Theta(\log N)$ contact nodes (called *fingers*), where N is the network size, and the i -th finger of X is the owner (or successor) of the ID $id_X + 2^{i-1}$ (the first finger, i.e., $i = 1$, is X 's direct successor). Besides, each node maintains a list of successor nodes for stabilization. Some DHTs [24] also utilize the list of successors during lookups to speed up the lookup process in the last few hops.

We study the more general case where the successor lists are used in lookups. We refer to the combination of the fingertable and the successor list as the *routing table*. More specifically, we consider the following lookup procedure. Assuming a lookup initiator I wants to find the owner of value v , it first queries n_1 , the node that is closest to v in its routing table, and n_1 sends its routing table to I .² Then out of n_1 's routing table, I finds the node n_2 closest to v and asks n_2 for its routing table. The process iteratively proceeds until reaching n_k , for which one of its successors is the owner of v (i.e., the lookup target).

Some of the queried nodes could be malicious and attempt to launch the following attacks.

1) *Lookup Bias Attack*: If the last queried node n_k is malicious, it can replace the honest nodes in its successor list with malicious nodes, so that one of its (malicious) “successors” will be concerned as the lookup target.

2) *Lookup Misdirection Attack*: Instead of trying to bias the lookup result, malicious nodes could attempt to make I query more malicious nodes during the lookup, by providing manipulated fingertables (i.e., replacing honest fingers with malicious nodes). This attack is a big threat to anonymity since the adversary can learn more information about the lookup target from a larger number of queried malicious nodes [14].

3) *Finger Pollution Attack*: In a DHT system, each node periodically performs lookups to update its fingers. An attack related to this is that malicious nodes could attempt to pollute honest nodes' fingertables during the finger-update lookups, so that the polluted fingertables can contribute to the lookup bias and misdirection attacks.

²In vanilla DHT lookups, I tells v to each queried node, which will return the finger closest to v ; however, this reveals the lookup target to malicious intermediate nodes. Hence, for anonymous lookups, I asks each intermediate node for its full routing table, without revealing v .

B. Security Mechanisms

Many existing secure DHT designs [6], [15], [16] employ redundant queries or lookups to tolerate misinformation provided by malicious nodes. However, the redundant transmission creates more opportunities for an adversary to gain information about the lookup initiator and/or target [13]. Some schemes [17], [18] utilize quorum-based topologies and threshold cryptography to limit Byzantine adversaries, but they also require the initiator to contact multiple nodes at each step of the lookup (for cryptographic operations), which accelerates information leaks. Myrmic [25] prevents routing table manipulation by introducing a central trusted authority to sign each node's routing table; however, this approach is impractical since for each node join or churn, the authority has to regenerate the signatures for all related nodes, rendering a performance bottleneck.

To effectively limit active attacks while minimizing information leaks, we propose a new defense strategy by letting each (honest) node secretly check the correctness of other nodes' routing tables. Such checks are performed offline (i.e., independent of lookups), and thus do not reveal any information of lookups. To punish discovered malicious nodes, we use a certificate authority (CA) to issue certificates and revoke certificates from identified malicious nodes so that the malicious nodes can be gradually removed from the system. We note that the CA in our case is fundamentally different from that of Myrmic [25]. The latter is required to be online all the time and needs to update signatures for multiple nodes for each node churn/join. Whereas, the certificates in our scheme are independent of nodes' routing states and thus do not need to be updated frequently. Our simulation results show that the workload of our CA is sufficiently low and can be handled by a modestly powerful server.

There has been several fairly efficient and scalable revocation mechanisms in the literature, such as Merkle Hash Tree based certificate revocation [26], efficient distribution of revocation information over P2P networks [27], and scalable PKI based on P2P systems [28]. Since certificate management in our scheme is essentially the same as these systems, we do not specifically study certificate revocation in this work.

1) *Secret Neighbor Surveillance*: To limit the lookup bias attack, we propose secret neighbor surveillance, a mechanism that prevent malicious neighbors from manipulating their successor lists.

In particular, we let each node maintain a predecessor list, in the same way as maintaining the successor list (i.e., periodically running Chord stabilization protocol anti-clockwise). The predecessor list is of the same size as the successor list, and thus each node X should be contained in the successor list of any of its predecessors. In other words, if X is not contained in the successor list of its predecessor, it

means this predecessor is trying to manipulate its successor list by replacing X with another node. Our goal is to let X detect this.




Figure 1: Secret neighbor surveillance. X is checking if itself is included in its predecessor P_3 's successor list. If not, it means P_3 manipulates its successor list by replacing X with another node.

In the lookup bias attack, a malicious node provides a manipulated successor list in response to lookup queries. Therefore, we let X *anonymously* send a "lookup query" to one of its predecessors, say P_3 (see Figure 1), and checks if itself is included in P_3 's successor list. Anonymity is required since if the malicious node can distinguish a testing query from real lookup queries based on the querier's identity, it can always provide the correct successor list for testing queries to avoid being detected. The anonymous transmission can be achieved by using the basic onion routing technique [29], i.e., X chooses two randomly peers as relays to forward its query to P_3 while using onion encryption to ensure each hop on the forwarding path can only know its previous and next hops. The two relay nodes can be found by performing a l -hop random walk on the overlay network (where $l = \Theta(\log N)$). Due to space limitation, we refer interested readers to our technical report [30] for the details of the random walk process.

X performs the above checks from time to time (i.e., with time interval $t_c \in_R (0, T_m]$ where T_m is the maximum checking interval) on randomly selected successors. A detected malicious node will be reported to the CA. To provide a non-repudiation proof on a manipulated successor list (i.e., verifiable to the CA), we let each node sign its routing table and attach a time stamp to it.

Another strategy to launch the lookup bias attack is to pollute honest nodes' successor lists during stabilization. For example, as shown in Figure 2, assuming P_2 is malicious and P_3 is honest, P_2 can send P_3 a manipulated successor list excluding X during P_3 's stabilization, so that P_3 will concern X as dead and remove X from its successor list; consequently, P_3 will be mistakenly identified as a malicious successor by X and P_2 will still be uncovered. To deal with this, we let each node sign its successor list used in stabilization; also, each node keeps a queue of latest received successor lists in stabilization as *proof*, to prove that its successor list is not intentionally manipulated. For

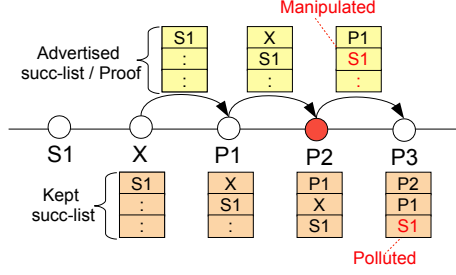


Figure 2: Successor list pollution. The malicious successor P_2 pollutes P_3 's successor list in P_3 's stabilization by providing a manipulated successor list.

example, P_3 can provide its proof to the CA showing that its successor list is correctly computed according to the information provided by P_2 . If P_3 's proof is verified (according to the stabilization algorithm) by the CA, then the suspicion on P_3 is cleared and the CA will request P_2 for its proof, and check it against P_3 's proof. This process is repeated until finding a node that cannot provide valid proofs and this node is then judged as the malicious node.

2) *Secret Finger Surveillance*: Likewise, we propose a *secret finger surveillance* mechanism for the lookup misdirection attack by limiting malicious nodes from manipulating their fingertables in the lookup.




Figure 3: Secret finger surveillance. X is checking if node Y has replaced its finger F with a malicious node F' . X first asks F' for its predecessor list, and then anonymously sends a “lookup query” to a random node P_1 in F' 's predecessor list. If any node in P_1 's successor list is closer to the ideal finger ID than F' , X detects Y manipulated its fingertable.

In particular, we let each node keep a small number of received fingertables (e.g., from lookups, secret neighbor surveillance, or random walks). From time to time, node X chooses a random finger from one of the kept fingertables, say the i -th finger F' of node Y , and asks F' for its predecessor list (see Figure 3). Then, after waiting a short random period of time, X anonymously sends a “lookup query” to a random predecessor of F' (say P'_1), and checks if any node in P'_1 's successor list is closer to the ideal finger ID than F' (i.e., F' is not the true finger F).

The intuition behind this is that if Y replaces a honest finger F with a malicious node F' , at least one of F' 's (true)

predecessors should be closer to the ideal finger ID than F' . Hence, if F' provides X with its true predecessor list, the fingertable manipulation will be detected. Therefore, F' has to manipulate its predecessor list to ensure that all the “predecessors” are malicious, so that the selected predecessor P'_1 can collude with F' by providing a manipulated successor list that is consistent with the predecessor list provided by F' . On the other hand, however, P'_1 cannot freely manipulate its successor list, since P'_1 is under surveillance by its neighbors (i.e., secret neighbor surveillance). Therefore, if the adversary tries to manipulate a single finger ($F \rightarrow F'$), she has to sacrifice at least one malicious node, either P'_1 or F' and Y .

3) *Secure Finger Update*: We can invoke the secret finger surveillance to limit the finger pollution attack: when X obtains the result (say F') of the finger-update lookup, it asks F' for its predecessor list, and chooses a random predecessor P'_1 of F' to perform the same checks as in the secret finger surveillance to verify F' is the true finger; X uses F' to update its fingertable only when F' passes these checks.

C. Security Evaluation

We use the following metrics to evaluate our security mechanisms.

- *fraction of remaining malicious nodes*,
- *false positive rate*, i.e., the chance that a honest node is judged as a malicious node,
- *false negative rate*, i.e., the chance that a malicious node is not identified when being tested by a node,
- *false alarm rate*, i.e., the chance that there is no node identified in a report sent to the CA.

These metrics represent different aspects of security properties. Reduction of malicious nodes shows *effectiveness*, false positive/negative rates represent *accuracy*, and false alarm rate indicates *efficiency*.

1) *Experiment Setup*: We developed an event-based simulator in C++ with about 3.0 KLOC. We consider a WAN setting, where latencies between each pair of peers are estimated using the King dataset [31]. We model node churn/join as an exponential distribution process $f(x) = \lambda e^{-(1/\lambda)x}$ with mean life time λ minutes. We generate random network topologies of size $N = 1000$ with 20% malicious nodes. Each node maintains 12 fingers and 6 successors/predecessors (on the order of $\Theta(\log N)$). With similar configurations as related work [1], [9], we let each node run successor and predecessor stabilization protocols every 2 seconds, and performs finger update every 30 seconds. To discover malicious nodes, each peer performs security checks of secret neighbor surveillance and secret finger surveillance every 60 seconds³. To ensure high identification

³This is based on the frequency of stabilization and finger update as well as the node churn rate, and we found that doing security checks every 60s is sufficient to rapidly discover malicious nodes.




Figure 4: Simulation results of the security mechanisms

accuracy, each node keeps 6 latest received successor lists as proofs. We let each node perform one lookup every minute (we choose 1 min only because we want to test a large number of lookups within a relatively short simulation time).

2) *Experimental Results:* We see from Figure 4a that the secret neighbor surveillance mechanism can rapidly identify malicious nodes that try to bias lookups. After a short time (20 mins), almost all malicious nodes are discovered. In comparison, the speed of discovering malicious nodes by the secret finger surveillance mechanism is relatively slower (as shown in Figure 4b), but still it can identify over 80% malicious nodes within 30 mins. The speed of identifying malicious nodes by the secure finger update mechanism is faster than that of the secure finger surveillance (as shown in Figure 4c), because the former is performed more frequently (at each finger update) and some malicious fingers contained in the successor list can also be detected by the secret neighbor surveillance.

The accuracy of our attacker discovery mechanisms is shown in Table I. The false positive rate is 0 for all the three mechanisms even when the churn rate is very high (e.g., the mean life time for each node is 10 mins). This ensures that honest nodes will not be judged as malicious nodes by mistake. In addition, the secret neighbor surveillance has very low false negative rate (less than 0.6%), which implies that any malicious nodes that try to bias lookups can be caught with high probability. We also see that the false negative rates for the secret finger surveillance and secure finger update mechanisms are relatively higher. This is because a malicious finger can pass the security checks if the randomly selected predecessor happens to be a colluding node and provides a successor list consistent with the malicious finger. However, over time, a malicious node can be identified with very high probability as shown in Figure 4.

We also compare our scheme with a state-of-the-art secure DHT scheme Halo [16] in terms of the number of biased lookups over time. We calculate the ratio of biased lookups of Halo according to their analysis results [16] §4.1 using the parameter $l = 7$ as they suggested. We can see from Figure 5 that after a short period of time, there are no more biased

lookups in Octopus, while the number of biased lookups of Halo keeps increasing linearly with the total number of lookups. This demonstrates that our security mechanisms can fundamentally thwart active attacks.




Figure 5: Security comparison

Finally, we evaluate the workload of the CA in terms of the number of messages (including reports, proofs, and etc) processed over time. We can see from Figure 6 that even during the peak time (the first 10 min), the CA only needs to process about 2 messages per second on average.




Figure 6: The CA's workload

IV. ANONYMITY MECHANISMS OF OCTOPUS

A. Problem Description

In vanilla DHT systems, since the lookup initiator I queries intermediate nodes directly, the queried nodes can easily infer the I 's identity. A natural idea to let I hide its identity by sending queries through an anonymous path, as in attacker identification mechanisms. An illustration of this is shown in Figure 7.

Table I: False positive/negative/alarm rates of the security mechanisms. λ is mean life time of each node (in minute). Attack rate is 100%. In fingertable manipulation/pollution attacks, checked malicious predecessors provide manipulated successor lists with 50% chance.

Security mechanisms	False Positive		False Negative		False Alarm	
	$\lambda = 60m$	$\lambda = 10m$	$\lambda = 60m$	$\lambda = 10m$	$\lambda = 60m$	$\lambda = 10m$
Secret neighbor surveillance	0	0	0	0.52%	0	0.52%
Secret finger surveillance	0	0	14.02%	19.55%	0.18%	1.55%
Secure finger update	0	0	14.08%	18.48%	0.33%	2.18%




Figure 7: T is the lookup target, $I \rightarrow A \rightarrow B$ is the anonymous path, and E_i 's are queried nodes.

However, we note that *a single anonymous path is insufficient to achieve high levels of anonymity*. We use the example in Figure 7 to show this. Assume queried nodes E_2 and E_4 are malicious, and the first relay A is also malicious. With a single anonymous path, the adversary can learn that E_2 and E_4 belong to the same lookup since they are contacted by the same exit node B . Wang et al. [14] have shown that based on the positions of a few queried malicious nodes in the lookup, the adversary can narrow the range of the lookup target into a small set of nodes (called *range estimation attack*). Suppose there are c concurrent lookups each having an estimation range of d nodes; then, the adversary can know that I is doing a lookup and its target is one of the $c \cdot d$ nodes.

B. Anonymity Mechanisms

To address the limitation of a single anonymous path, we propose to split lookup queries over multiple anonymous paths, as shown in Figure 8.




Figure 8: The structure of multiple anonymous paths in Octopus. A , B , C_i and D_i are relays.

Using separate anonymous paths for different queries effectively disassociates the adversary's observations. The

adversary only sees disjoint events from different concurrent lookups, but is unable to group queries belonging to the same lookup; in this case, it is much harder to apply the range estimation attack, thus substantially limiting the information leak.

Moreover, to further blur the adversary's observations and make the range estimation attack even harder, we propose to add *dummy queries* in the lookup. Then, even though in rare cases the adversary can link two queried nodes in the same lookup (e.g., when all the relays used for the two queries are malicious), the adversary is unable to tell whether they are dummy queries or true queries, and the result of the range estimation attack would be incorrect with a dummy query.

We note that using multiple anonymous paths is important to ensure effectiveness of dummy queries, because in the single-anonymous-path scenario, observed queries are linkable due to the common exit relay and hence dummy queries can be distinguished based on the positions of observed queries. In comparison, with multiple anonymous paths, identifying dummy queries is much harder.

C. Anonymity Evaluation

We analyze the best strategies for an adversary to infer the lookup target T and the initiator I based its observations, and calculate the target anonymity $H(T)$ and the initiator anonymity $H(I)$. We use entropy to quantify $H(T)$ and $H(I)$. We let O denote the set of possible observations of the adversary (including null observation). To measure the system as a whole, we have:

$$H(T) = \sum_{o \in O} P(o) \cdot H(T|o), H(I) = \sum_{o \in O} P(o) \cdot H(I|o) \quad (1)$$

where $P(o)$ is the probability of observation o occurring.

To calculate the maximum information leak, we make the following assumptions in the anonymity calculation. First, we assume the network is static, since network dynamics can obscure the adversary's observations and make it more difficult to extract information about the initiator/target. Second, we only consider passive attacks, as active attackers can be quickly identified by our security mechanisms and consequently the adversary will lose observers to carry out passive attacks.

D. Anonymity Calculation

Due to space limitation, we only provide the outline of calculating the initiator anonymity. We refer interested readers to [30] for detailed anonymity calculation and analysis.




Figure 9: Anonymity evaluation of Octopus. a is the concurrent lookup rate.




Figure 10: Anonymity comparison. $a = 1\%$.

To calculate the initiator anonymity, we divide the observations of the adversary into two categories.

- o_n : the observation occurring when T is not observed
- O_o : the set of observations occurring when T is observed

According to 1, $H(I)$ is calculated as:

$$H(I) = P(o_n) \cdot H(I|o_n) + \sum_{o_o \in O_o} P(o_o) \cdot H(I|o_o) \quad (2)$$

When T is not observed, the entropy of I is maximized. Presuming that the adversary can exclude malicious nodes from the anonymity set of I , we have:

$$H(I|o_n) = \log_2((1-f) \cdot N) \quad (3)$$

Let \mathcal{R}_T^l denote the set of non-dummy queries linkable to I in the lookup whose target is T . Then, based on whether \mathcal{R}_T^l is an empty set, we calculate $H(I|o_o)$ as follows:

$$H(I|o_o) = P(\mathcal{R}_T^l = \emptyset) \cdot H'(I|o_o) + (1 - P(\mathcal{R}_T^l = \emptyset)) \cdot H''(I|o_o) \quad (4)$$

When there is no linkable non-dummy query, I is un-linkable with T . However, since some of the initiators of concurrent lookups can be observed by the adversary, we have:

$$H'(I|o_o) = P(I_{obsv}) \cdot \log_2(\#obsv_{hon_init}) + (1 - P(I_{obsv})) \cdot \log_2((1-f) \cdot N) \quad (5)$$

Let Ψ^l denote the set of concurrent lookups that have at least one linkable query, and ψ_T denote the lookup with

target T . When $\mathcal{R}_T^l \neq \emptyset$, $\psi_T \in \Psi^l$. Each lookup in Ψ^l is possible to be ψ_T . Therefore, we have:

$$H''(I|o_o) = - \sum_{\psi \in \Psi^l} P(\psi = \psi_T|o_o) \cdot \log_2 P(\psi = \psi_T|o_o) \quad (6)$$

Because the density of queries close to the target is higher than other regions on the ring, for ψ_T it is highly likely that the last queried node in \mathcal{R}_T^l is located very close to T . Therefore, the adversary can assign probability to each candidate initiator based on the minimum distance (i.e., number of hops) between its queried nodes and T . In particular, let \mathcal{Q}_ψ^l denote the set of linkable queries in ψ , $\psi \in \Psi^l$, and let $\xi(x)$ denote the probability that for ψ_T the minimum distance from linkable queried nodes to T is x . $\xi(x)$ can be obtained via pre-simulations of the lookup. Then, we can calculate $P(\psi = \psi_T|o_o)$ as follows:

$$P(\psi = \psi_T|o_o) \approx \frac{\xi(\min_{E \in \mathcal{Q}_\psi^l} \text{dist}(E, T))}{\sum_{\psi' \in \Psi^l} \xi(\min_{E' \in \mathcal{Q}_{\psi'}^l} \text{dist}(E', T))} \quad (7)$$

E. Results and Comparisons

We developed a simulator for anonymity measurements in C++ with about 1.3 KLOC. The results are shown in Figure 10. With network size $N = 100\,000$, concurrent lookup rate $a = 1\%$, $f = 20\%$ malicious nodes, and 6 dummies, Octopus only leaks 0.57 bits of information about the initiator and 0.82 bits of information about the target. We compare Octopus with the base-line scheme Chord [1] and

Table II: Performance comparison. τ is the time interval between two consecutive lookups.

Schemes	Lookup Latency (sec)		Bandwidth Consumption (Kbps)	
	Mean	Median	$\tau = 5min$	$\tau = 10min$
Octopus	2.15	1.61	5.91	4.30
Chord [1]	1.35	0.35	0.29	0.28
Halo [16]	6.89	1.79	0.71	0.37

the state-of-the-art anonymous DHT lookups NISAN [7] and Torsk [8] (we do not explicitly compare Octopus with secure DHTs such as Halo [16], since they leak more information than Chord.). We can see that in the same setting, NISAN and Torsk leak about 3.3 bits of information about the initiator, which is about 6 times more than Octopus. As for the target anonymity, the information leak for NISAN and Torsk is 11.3 bits and 3.4 bits, which is 13 times and 4 times more than that of Octopus, respectively.

V. PERFORMANCE EVALUATION

A. Lookup Latency

Lookup latency is one of the most important performance factors for DHT systems. We measure the lookup latency of Octopus using PlanetLab with 207 randomly selected nodes. We use boost C++ library⁴ (mainly UDP asynchronous read/write of Boost.Asio) to build the communication substrate. We let each node perform 2000 lookups independently using randomly picked lookup keys. For each lookup, we record the latency from the time of sending out the first query till the time of receiving the lookup result.




Figure 11: Comparison of lookup latency on Planetlab.

For comparison, we use the same methodology to implement Chord [1] and Halo [16], and measure their lookup latencies in the same network environment. We choose Halo because it is one of the state-of-the-art secure DHT lookup schemes and it is also based on Chord overlay. For Halo, we use degree-2 recursion with redundant parameter 8×4 , as suggested by Kapadia and Triandopoulos [16] to provide fairly strong security guarantee. The experimental results are presented in Figure 11 and Table II. We can see that while the lookup latency of Octopus is relatively longer than that of Chord due to more transmission for security and anonymity needs, it is smaller than that of Halo, which only provides

security guarantees. The performance advantage is because Octopus does not rely on redundant lookups, while in Halo a lookup is not completed until all redundant lookups' results are returned.

B. Bandwidth Overhead

We also compare Octopus with Chord and Halo in terms of bandwidth cost. We adopt the same configuration as described in Section III-C1 for each of the DHT lookups, and consider an overlay network with 1 000 000 nodes⁵. We can see from Table II that Octopus does incur higher communication overhead than Chord and Halo, in order to achieve high levels of anonymity; however, the bandwidth cost of Octopus is still reasonable (only a few Kbps), which is affordable even for low-end clients with limited bandwidth.

VI. RELATED WORK

A. Secure DHT Lookups

A major school of proposals to securing DHT lookups uses *redundancy*. Castro et al. [15] proposed a robust DHT system that relies on redundant lookups. Each key is replicated among several replica nodes (typically the neighbors of the key owner). Instead of doing a single lookup, the initiator performs multiple redundant lookups towards all the replicas. The lookup result would be correct as long as one of the redundant lookups is not biased. The limitation of this approach is that the redundant lookups tend to converge to a small number of nodes close to the target, and one malicious node in this set could infect many redundant lookups. Much subsequent work (such as [6], [16], [32]) focuses on disentangling the redundant lookup paths to provide better security. Cyclone [32] partitions nodes into r Chord sub-rings based on similarity of node IDs, and has r redundant lookups routed through the r sub-rings independently. Salsa [6] uses a new virtual-tree-based DHT structure, in which any two nodes share few global contacts so that redundant messages can proceed along different paths. Halo [16] does not change the underlying DHT structures, but uses the original Chord overlay and performs redundant searches towards *knuckles* — nodes that have fingers pointing to the target.

While effective in ensuring security, these redundant-lookup-based approaches are incapable of preserving anonymity, since redundant transmission creates opportunities for an adversary to gain information about the lookup initiator and/or target [13]. ShadowWalker [9] embeds redundancy into the DHT itself and uses *shadows* (nodes in

⁵We use the following parameters to estimate the bandwidth overhead. Each routing state item (such as fingers or successors) is 10 bytes. We use ECDSA signature (40 bytes) for authentication with a 4-byte timestamp, and AES-128 for onion encryption. Each certificate is 50 bytes, including the node's IP address (6 bytes), the node's public key (20 bytes), expire time (4 bytes), and the CA's signature (20 bytes).

⁴www.boost.org

redundant topologies) to verify each step of a lookup. Unfortunately, Schuchard et al. [33] found that ShadowWalker is vulnerable to *eclipse attack*, where the entire set of shadows of a certain node are compromised, leading to other nodes' routing states being infected. They also showed that increasing the dimension of redundant topologies can mitigate the eclipse attack, but at a higher performance cost.

Another major school of research on secure DHT lookups leverages cryptographic techniques. Myrmic [25] uses an online certificate authority to sign each node's routing state. The major limitation of Myrmic is that for each node join/churn, the central authority has to update the certificates for all related nodes. Young et al. [17] proposed two schemes RCP-I and RCP-II that use threshold signature and distributed key generation to avoid the reliance of a central authority. In their schemes, the verification information on each message is collaboratively generated by a threshold number of nodes, rather than by a central authority.

All these secure DHT lookup schemes are not designed to preserve anonymity. Lookup keys are revealed during queries, and identities of lookup initiators are easily exposed due to directly contacting intermediate nodes.

B. Secure and Anonymous DHT Lookups

NISAN [7] is among the first to try to provide both security and anonymity guarantees in DHT systems. For security purpose, each queried node is required to provide its entire fingertable, so that the lookup initiator can apply *bound checking* on it to limit manipulation of fingertables. NISAN also uses redundancy to enhance security. The authors proposed a greedy search mechanism to query multiple nodes at each step and combine the query results to tolerate misinformation. On the other hand, acquiring the entire fingertable also helps protect the anonymity of lookup targets, since the lookup keys are not revealed to intermediate nodes. Nevertheless, NISAN can only provide very limited anonymity protection. Wang et al. [14] showed that a passive adversary is able to narrow the range of a lookup target down to a small number of nodes, by analyzing the locations of observed queries (called a *range estimation attack*).

Torsk [8] is a DHT-based anonymous communication system. A key component of Torsk is a proxy-based anonymous DHT lookup. The idea is that a lookup initiator performs a random walk on the overlay to find a random node (called *buddy*), and requests the buddy to perform the lookup on its behalf. Because Torsk uses Myrmic [25] to secure lookups, it has the same limitation as Myrmic — requiring an online central authority to sign each node's routing state. In addition, as we analyzed in Section IV-A, a single proxy structure is insufficient to provide high levels of anonymity: the information learned by the range estimation attack can be used to launch *relay exhaustion attack* [14].

Recently, Backes et al. [18] proposed to leverage *oblivious transfer* to add query privacy to RCP-I and RCP-II [17]. However, for similar reasons as NISAN, this scheme is vulnerable to the range estimation attack, since the initiator needs to contact multiple intermediate nodes at each step of the lookup.

Freenet [34] is a deployed P2P system, which allows people to upload sensitive files to the overlay and employs data duplication strategies to make them hard to block. Freenet aims to preserve the publishers' privacy, but does not provide anonymity in lookups. Vasserman et al. [35] create a membership concealing overlay network (MCON) for unobservable communication. They aim to make it difficult for either an insider or outsider adversary to learn the set of participating members. This is similar to previous *darknet* designs [36]. However, MCON and darknets are not designed to provide anonymity.

VII. CONCLUSION

In this paper, we presented Octopus, a new DHT lookup that provides strong guarantees for both anonymity and security. Octopus ensures security and anonymity via three fundamental techniques. First, Octopus constructs an anonymous path to send lookup query messages while hiding the initiator. Second, it splits the individual queries used in a lookup over multiple paths, and introduces dummy queries, to make it difficult for an adversary to learn the lookup target. Third, it uses secret security checks to identify and remove malicious nodes. We developed an event-based simulator, and showed that malicious nodes can be quickly identified with high accuracy. In addition, via probabilistic modeling and simulation, we showed that Octopus can achieve near-optimal anonymity for both the lookup initiator and target. We also evaluated the efficiency of Octopus on Planetlab, and showed that Octopus has reasonable lookup latency and communication overhead.

VIII. ACKNOWLEDGEMENTS

We would like to thank Prateek Mittal for helpful discussion throughout the project. We also thank George Danezis for invaluable feedback on the earlier draft of the paper. In addition, we are grateful to Prateek Mittal for sharing his codes of an event-based simulator and a DHT implementation on Planetlab, which we used as a reference for our implementations. Finally, we thank anonymous reviewers for helpful comments to improve the paper. This research was supported in part by NSF CNS 09-53655.

REFERENCES

- [1] I. Stoica, R. Morris, D. Liben-Nowell, D. R. Karger, M. F. Kaashoek, F. Dabek, and H. Balakrishnan, "Chord: A scalable peer-to-peer lookup protocol for internet applications," *IEEE/ACM Trans. Netw.*, vol. 11, no. 1, pp. 17–32, 2003.
- [2] P. Maymounkov and D. Mazieres, "Kademlia: A peer-to-peer information system based on the xor metric," *IPTPS*, 2001.

- [3] M. J. Freedman, E. Freudenthal, and D. Mazières, “Democratizing content publication with coral,” in *Proc. of NSDI’04*, 2004.
- [4] A. Rowstron and P. Druschel, “Storage management and caching in past, a large-scale, persistent peer-to-peer storage utility,” in *Proc. SOSP’01*, 2001.
- [5] A. Mislove, G. Oberoi, A. Post, C. Reis, P. Druschel, and D. S. Wallach, “Ap3: Cooperative, decentralized anonymous communication,” *ACM SIGOPS European Workshop*, 2004.
- [6] A. Nambiar and M. Wright, “Salsa: A structured approach to large-scale anonymity,” *ACM CCS*, 2006.
- [7] A. Panchenko, S. Richter, and A. Rache, “Nisan: Network information service for anonymization networks,” *ACM CCS*, November 2009.
- [8] J. McLachlan, A. Tran, N. Hopper, and Y. Kim, “Scalable onion routing with tors,” *ACM CCS*, November 2009.
- [9] P. Mittal and N. Borisov, “Shadowwalker: Peer-to-peer anonymous communication using redundant structured topologies,” *ACM CCS*, November 2009.
- [10] A. Shakimov, A. Varshavsky, L. P. Cox, and R. Cáceres, “Privacy, cost, and availability tradeoffs in decentralized osns,” ser. WOSN ’09, 2009, pp. 13–18.
- [11] L. A. Cuttillo, R. Molva, and T. Strufe, “Privacy preserving social networking through decentralization,” *Proceedings of the Sixth international conference on Wireless On-Demand Network Systems and Services*, 2009.
- [12] D. Wallach, “A survey of peer-to-peer security issues,” in *Software Security Theories and Systems*, ser. LNCS, 2003, vol. 2609, pp. 253–258.
- [13] P. Mittal and N. Borisov, “Information leaks in structured peer-to-peer anonymous communication systems,” *ACM CCS*, 2008.
- [14] Q. Wang, P. Mittal, and N. Borisov, “In search of an anonymous and secure lookup: Attacks on structured peer-to-peer anonymous communication systems,” *ACM CCS*, 2010.
- [15] M. Castro, P. Druschel, A. Ganesh, A. Rowstron, and D. S. Wallach, “Secure routing for structured peer-to-peer overlay networks,” in *OSDI*, December 2002.
- [16] A. Kapadia and N. Triandopoulos, “Halo: High-assurance locate for distributed hash tables,” in *NDSS*, February 2008.
- [17] M. Young, A. Kate, I. Goldberg, and M. Karsten, “Practical robust communication in dhts tolerating a byzantine adversary,” *Proc. ICDCS’10*, pp. 263–272, 2010.
- [18] M. Backes, I. Goldberg, A. Kate, and T. Toft, “Adding query privacy to robust dhts,” *Tech. rep., arXiv:1107.1072v1 [cs.CR]*, July 2011.
- [19] J. Douceur, “The sybil attack,” in *Peer-to-Peer Systems*, 2002, vol. 2429, pp. 251–260.
- [20] N. Borisov, “Computational puzzles as sybil defenses,” in *Peer-to-Peer Computing*, 2006.
- [21] G. Danezis and P. Mittal, “Sybilinifer: Detecting sybil nodes using social networks,” in *NDSS*, 2009.
- [22] H. Yu, P. B. Gibbons, M. Kaminsky, and F. Xiao, “Sybillimit: A near-optimal social network defense against sybil attacks,” in *IEEE Symposium on Security and Privacy (Oakland)*, 2008.
- [23] A. Pfizmann and M. Hansen, “A terminology for talking about privacy by data minimization: Anonymity, unlinkability, undetectability, unobservability, pseudonymity, and identity management,” Aug. 2010, v0.34.
- [24] A. Rowstron and P. Druschel, “Pastry: Scalable, decentralized object location and routing for large-scale peer-to-peer systems,” *IFIP/ACM International Conference on Distributed Systems Platforms (Middleware)*, pp. 329–350.
- [25] P. Wang, I. Osipkov, N. Hopper, and Y. Kim, “Myrmic: Secure and robust dht routing,” *Tech. rep., Digital Technology Center, University of Minnesota at Twin Cities*, 2007.
- [26] J. L. Muoz, J. Forne, O. Esparza, and M. Soriano, “Certificate revocation system implementation based on the merkle hash tree,” *Inte. Journ. of Inf. Sec.*, vol. 2, no. 2, 2003.
- [27] M. C. Morogan and S. Muftic, “Certificate revocation system based on peer-to-peer crt distribution,” in *Proc. of the DMS’03 Conference*, 2003.
- [28] C. Y. Liau, S. Bressan, K.-L. Tan, C. Yee, L. Stphane, and B. K. lee Tan, “Efficient certificate revocation: A p2p approach,” in *HICSS’05*, 2005.
- [29] P. Syverson, G. Tsudik, M. Reed, and C. Landwehr, “Towards an analysis of onion routing security,” *International Workshop on Design Issues in Anonymity and Unobservability*, vol. 2009, LNCS, Springer, July 2000.
- [30] Q. Wang and N. Borisov, “Octopus: A secure and anonymous dht lookup,” technical report, arXiv cs.CR, 2012. <http://arxiv.org/abs/1203.2668>.
- [31] F. Dabek, J. Li, E. Sit, J. Robertson, M. F. Kaashoek, and R. Morris, “Designing a dht for low latency and high throughput,” *NSDI’04*, 2004.
- [32] M. S. Artigas, P. G. Lopez, J. P. Ahullo, and A. F. G. Skarmeta, “Cyclone: A novel design schema for hierarchical dhts,” *In P2P 05*, pp. 49–56, 2005.
- [33] M. Schuchard, A. Dean, V. Heorhiadi, N. Hopper, and Y. Kim, “Balancing the shadows,” *ACM WPES*, 2010.
- [34] I. Clarke, T. W. Hong, S. G. Miller, O. Sandberg, and B. Wiley, “Protecting Free Expression Online with {Freenet},” *IEEE Internet Computing*, vol. 6, no. 1, pp. 40–49, 2002.
- [35] E. Y. Vasserman, R. Jansen, J. Tyra, N. Hopper, and Y. Kim, “Membership-concealing overlay networks,” in *ACM CCS*, 2009.
- [36] WASTE. <http://waste.sourceforge.net/>.